PROJECT HEALTH PULSE

High-Level Problem Statements & Strategic Asks

A client-ready discussion document focused on improving project intelligence, early warning, operational efficiency, explainability, and AI cost visibility.

Discussion objective	Align on the highest-value enhancements and identify the capabilities required to move Project Health Pulse from project reporting toward predictive, actionable and economically measurable project intelligence.

The asks below are intentionally framed as discussion points rather than prescriptive implementation decisions.

1. High-Level Problem Statements / Asks

1. Manual & reactive project health assessment

Ask: How can we move from manually reviewing meeting discussions to an AI-driven, evidence-based project health assessment that identifies risks early?

2. Long processing time

Ask: How can we reduce the time taken after transcript ingestion by parallelizing processing, optimizing model usage, and reusing previously processed information?

3. AI cost is difficult to control

Ask: How can we get visibility into the end-to-end cost of an AI assessment, instead of looking only at individual LLM/token costs?

4. Lack of business transaction-level costing

Ask: How can we associate every AI and non-AI cost with a single business transaction or assessment ID so that we can calculate the actual cost per business outcome?

5. Limited evidence and explainability

Ask: How can we make every project-health rating traceable back to the exact source, meeting, and evidence that influenced the decision?

6. Reporting vs. prediction

Ask: How can we evolve the solution from reporting current project health to predicting future risks, milestone slippage, and potential escalations?

7. Recommendations are not closed-loop

Ask: How can we convert AI recommendations into owned actions and measure whether those actions actually improved project health?

8. Information is fragmented

Ask: How can we combine meeting transcripts with Jira, status reports, RAID logs, project plans, and other approved sources to create a more complete project picture?

9. No portfolio-level visibility

Ask: How can leadership see the health and risk of multiple projects from a single portfolio view and quickly identify where intervention is required?

2. Quick Ideas / Discussion Asks

Speed & Cost

- Can we parallelize the processing after Ingest?
- Can we use smaller models for simple tasks and stronger models only where reasoning is required?
- Can we cache previously processed transcripts and only process new information?
- Can we pre-process assessments in batch so login becomes mostly Load + Display?
- Can we calculate cost per assessment, per pillar, and per processing stage?

Trust & Evidence

- Can we show exact evidence behind every rating?
- Can we show what changed compared with the previous assessment?
- Can we allow feedback and corrections to improve the system?
- Can we maintain production-grade auditability for models, prompts, evidence, decisions, and user actions?

Functional Intelligence

- Can we automatically generate RAID items, decisions, and action items?
- Can we detect contradictions between different project meetings or sources?
- Can we add an Ask-the-Project chat interface?
- Can we add What-if analysis, such as the impact of adding an engineer?

Predictive Capabilities

- Can we predict milestone slippage before it happens?
- Can we identify early warning signals before the health score becomes Poor?
- Can we learn from historical project patterns?

Follow-through

- Can recommendations automatically become trackable actions with owners and due dates?
- Can we measure whether an AI recommendation actually improved the project?
- Can we automatically alert the right stakeholder when project health deteriorates?

Portfolio & Integration

- Can we create a portfolio-level RAG dashboard?
- Can we integrate approved sources such as Teams, Zoom, Jira, risk logs and status reports?
- Can we provide role-based views for delivery teams, leaders and clients?

AI FinOps / Transaction Economics

- Can we create one Assessment ID and track all AI and non-AI costs against it?
- Can we calculate true end-to-end cost per business transaction or assessment?

- Can we include abandoned, failed and reprocessed runs when measuring economics?
- Can we connect AI spend to the business outcome rather than stopping at token consumption?

3. Recommended Leadership Discussion

For an initial client discussion, the following five asks provide a focused narrative without overloading the conversation with implementation detail.

Theme	Leadership Ask
Predictive	Can we make Project Health Pulse predictive rather than just reporting?
Explainable	Can we make every AI decision evidence-backed and traceable?
Actionable	Can we automate RAID, actions, decisions, and follow-up?
Economically measurable	Can we calculate true end-to-end cost per business transaction or assessment?
Portfolio-ready	Can we provide portfolio-level visibility with early-warning alerts?

4. Intended Evolution

→ Current: project health reporting

→ Next: evidence-backed project intelligence

→ Next: automated RAID, actions and decisions

→ Next: trend analysis and early-warning prediction

→ Next: portfolio-level project intelligence

→ Strategic layer: end-to-end AI FinOps tied to business outcomes

North Star: Move from “What happened?” to “What is likely to happen, what should we do, and what does it cost to deliver that outcome?”

Prepared for client discussion • Discussion draft • Subject to scope, architecture and governance validation